

Assignment2A

ANS.1. • **R-squared vs. Residual Sum of Squares (RSS)**

R-squared is generally considered a better measure of the goodness of fit in a regression model compared to **Residual Sum of Squares (RSS)**. Here's why:

- **R-squared:** This statistic provides a measure of how well the regression model explains the variability of the dependent variable. It is the proportion of the variance in the dependent variable that is predictable from the independent variables. An R-squared value ranges from 0 to 1, where a higher value indicates a better fit, meaning that the model explains a larger portion of the variance.
- **Residual Sum of Squares (RSS):** This is the sum of the squared differences between the observed values and the values predicted by the model. While RSS provides information on the fit of the model, it is an absolute measure that depends on the scale of the data. Lower RSS values indicate a better fit, but it does not provide a relative measure of goodness of fit as R-squared does.

Therefore, R-squared is more commonly used as it provides a relative measure that is easier to interpret, showing the proportion of variance explained by the model.

• **Total Sum of Squares (TSS), Explained Sum of Squares (ESS), and Residual Sum of Squares (RSS)**

- **Total Sum of Squares (TSS):** This measures the total variance in the dependent variable. It is calculated as the sum of the squared differences between each observed value and the mean of the observed values. TSS represents the total variability in the data.

$$TSS = \sum_{i=1}^n (y_i - \bar{y})^2$$

where y_i are the observed values and \bar{y} is the mean of the observed values.

- **Explained Sum of Squares (ESS):** This measures the portion of the total variance that is explained by the regression model. It is the sum of the squared differences between the predicted values and the mean of the observed values.

$$ESS = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$$

where \hat{y}_i are the predicted values.

- **Residual Sum of Squares (RSS):** As mentioned, RSS measures the portion of the total variance that is not explained by the model. It is the sum of the squared differences between the observed values and the predicted values.

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

The relationship between these sums of squares is given by the equation:

$$TSS = ESS + RSS$$

This equation shows that the total variability (TSS) is partitioned into the variability explained by the model (ESS) and the unexplained variability (RSS). R-squared can be calculated using these sums of squares:

$$R^2 = \frac{ESS}{TSS} = 1 - \frac{RSS}{TSS}$$

This formula for R-squared illustrates that it represents the proportion of the total variance that is explained by the model.

ANS.2. TSS (Total Sum of Squares), ESS (Explained Sum of Squares), and RSS (Residual Sum of Squares) are important metrics used to assess the fit of a regression model. Here are their definitions and the equation that relates them:

1. Total Sum of Squares (TSS):

- TSS measures the total variance in the dependent variable. It represents the total variability in the observed data.
- Formula: $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ where y_i are the observed values and \bar{y} is the mean of the observed values.

2. Explained Sum of Squares (ESS):

- ESS measures the portion of the total variance that is explained by the regression model. It reflects how much of the variation in the dependent variable is accounted for by the independent variables.
- Formula: $ESS = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$ where \hat{y}_i are the predicted values from the regression model.

3. Residual Sum of Squares (RSS):

- RSS measures the portion of the total variance that is not explained by the model. It represents the variability in the dependent variable that the model fails to capture.
- Formula: $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$ where y_i are the observed values and \hat{y}_i are the predicted values.

The relationship between these three metrics is given by the equation:

$$TSS = ESS + RSS$$

This equation shows that the total variability in the data (TSS) is partitioned into the variability explained by the model (ESS) and the unexplained variability (RSS).

Additionally, R-squared (R^2), which is a measure of the goodness of fit, can be calculated using these sums of squares:

$$R^2 = \frac{ESS}{TSS} = 1 - \frac{RSS}{TSS}$$

ANS.3. R-squared represents the proportion of the total variance in the dependent variable that is explained by the independent variables in the model.

Regularization is a critical technique in machine learning to reduce overfitting, enhance model generalization, and manage model complexity. Several regularization techniques are used across different [types of models](#). Here are some of the most common and effective regularization techniques:

1. **L1 Regularization (Lasso):** Encourages sparsity in the model parameters. Some coefficients can shrink to zero, effectively performing feature selection.
2. **L2 Regularization (Ridge):** It shrinks the coefficients evenly but does not necessarily bring them to zero. It helps with multicollinearity and model stability.
3. **Elastic Net:** This is useful when there are correlations among features or to balance feature selection with coefficient shrinkage.
4. **Dropout:** Results in a network that is robust and less likely to overfit, as it has to learn more robust features from the data that aren't reliant on any small set of neurons.
5. **Early Stopping:** Prevents overfitting by not allowing the training to continue too long. It is a straightforward and often very effective form of regularization.
6. **Batch Normalization:** Reduces the need for other forms of regularization and can sometimes eliminate the need for dropout.
7. **Weight Constraint:** This constraint ensures that the weights do not grow too large, which can help prevent overfitting and improve the model's generalization.
8. **Data Augmentation:** Although not a direct form of regularization in a mathematical sense, it acts like one by artificially increasing the size of the training set, which helps the model generalize better.

Ans4. The Gini impurity index is a measure used in decision tree algorithms for classification tasks. It quantifies the impurity or uncertainty of a dataset, where a lower Gini impurity indicates purer nodes (more homogeneous class distribution) and a higher Gini impurity indicates impure nodes (more heterogeneous class distribution).

Here's a detailed explanation:

1. Definition:

- Gini impurity is a measure of how often a randomly chosen element from the set would be incorrectly labeled if it was randomly labeled according to the distribution of labels in the set.
- It is calculated by summing the probabilities of each item being chosen times the probability of a mistake in categorizing that item.

2. Formula:

- For a set S containing samples from K classes, the Gini impurity $Gini(S)$ is calculated as: $Gini(S) = 1 - \sum_{i=1}^K p_i^2$ where p_i is the probability of picking a sample of class i from the set S .

3. Interpretation:

- A Gini impurity of 0 indicates that the set S is completely pure, i.e., all elements belong to the same class.
- A Gini impurity of 0.5 indicates that the set S is completely impure, i.e., the elements are evenly distributed across all classes.

4. Decision Trees:

- In the context of decision trees, the Gini impurity is used to decide the order in which to split a dataset into subsets to maximize the information gain. The split that results in the lowest Gini impurity is chosen as the best split.
- The Gini impurity index is used in algorithms like CART (Classification and Regression Trees) to build decision trees by recursively splitting the dataset based on features that minimize the impurity of the resulting child nodes.

5. Comparison with Entropy:

- Gini impurity and entropy are two common measures of impurity used in decision trees. While Gini impurity is more efficient to compute, entropy tends to produce slightly more balanced trees.
- Gini impurity tends to isolate the most frequent class in its own branch of the tree, while entropy tends to produce slightly more balanced trees.

Ans5. Yes, unregularized decision trees are prone to overfitting. Here's why:

1. Flexibility and Complexity:

- Decision trees are non-parametric models that can fit complex decision boundaries. They are capable of learning intricate relationships between features and target variables.
- Without any constraints (regularization), decision trees can grow deep and complex, capturing noise and outliers in the training data.

2. Overfitting:

- Overfitting occurs when a model learns not only the underlying patterns in the training data but also noise and randomness. As a result, an overfitted model performs well on the training data but poorly on unseen test data.

- Unregularized decision trees have high variance. They may fit the training data very closely, resulting in high accuracy on training data but poor generalization to new data.
3. **High Variance:**
- Unregularized decision trees tend to have high variance because they can fit the training data very closely, which results in a model that is sensitive to small changes in the training data.

An6. An ensemble technique in machine learning is a method that combines multiple individual models (often called base learners or weak learners) to improve predictive performance. The basic idea behind ensemble methods is to leverage the diversity among models to achieve better results than any single model could alone. Ensemble methods are widely used and often result in more accurate predictions compared to individual models.

Here are some key points about ensemble techniques:

1. **Types of Ensemble Techniques:**

- **Bagging (Bootstrap Aggregating):** Involves training multiple models independently on different random subsets of the training data and then averaging their predictions. Example: Random Forest.
- **Boosting:** Trains multiple models sequentially, where each model tries to correct the errors of its predecessor. Examples: AdaBoost, Gradient Boosting Machines (GBM), XGBoost, LightGBM.
- **Stacking (Stacked Generalization):** Combines the predictions of multiple models as input features to a new model, which then learns to make final predictions.
- **Voting:** Combines predictions from multiple models either by taking the majority vote (for classification) or by averaging (for regression).

2. **Advantages:**

- **Improved Accuracy:** Ensemble methods often produce more accurate predictions compared to individual models.
- **Robustness:** They are less prone to overfitting, especially when combining diverse models.
- **Versatility:** Ensemble techniques can be applied to a variety of models and problems.

3. **Practical Considerations:**

- **Diversity:** The strength of an ensemble comes from the diversity of the individual models. They should make errors independently, so that they can complement each other when combined.
- **Computation:** Ensembles can be more computationally expensive and complex compared to single models, especially when combining multiple models.
- **Implementation:** Many machine learning libraries provide built-in support for ensemble methods, making them easier to implement.

4. **Examples:**

- **Random Forest:** An ensemble of decision trees, where each tree is trained on a random subset of the data and random subset of the features.

- **Gradient Boosting Machines (GBM):** An ensemble of decision trees, where each tree corrects errors made by the previous tree.
- **AdaBoost:** A boosting technique that combines multiple weak classifiers to create a strong classifier.

ans7. Bagging (Bootstrap Aggregating) and Boosting are two popular ensemble techniques in machine learning. While both methods involve combining multiple models to improve predictive performance, they have distinct approaches and characteristics. Here are the key differences between Bagging and Boosting:

Bagging (Bootstrap Aggregating):

1. **Approach:**
 - **Independence:** Bagging involves training multiple models independently on different random subsets of the training data.
 - **Parallelism:** Models are trained in parallel, and their predictions are averaged (or majority voted in case of classification) to make the final prediction.
2. **Training Process:**
 - **Bootstrap Sampling:** Each model is trained on a random sample of the training data with replacement (bootstrap sampling). This means that some samples may be repeated in the same subset, while others may not be included at all.
 - **Feature Randomization:** In addition to sampling data points, Bagging may also involve randomizing subsets of features for each model.
3. **Examples:**
 - **Random Forest:** A widely used Bagging ensemble of decision trees, where each tree is trained on a bootstrap sample of the data and random subset of features.
4. **Main Goal:**
 - **Reduce Variance:** Bagging aims to reduce variance, which helps to improve the stability and robustness of the model by reducing overfitting.

Boosting:

1. **Approach:**
 - **Sequential Learning:** Boosting involves training multiple models sequentially, where each model tries to correct the errors made by its predecessor.
 - **Adaptive Weighting:** Examples that were incorrectly predicted by the previous model are given more weight, so that subsequent models focus more on those examples.
2. **Training Process:**
 - **Weight Adjustment:** Each subsequent model pays more attention to instances that were incorrectly predicted by the previous model. This is done by adjusting weights or probabilities.
 - **Focus on Hard Examples:** Boosting focuses more on examples that are harder to classify.
3. **Examples:**

- **AdaBoost (Adaptive Boosting):** A popular boosting algorithm that combines multiple weak learners (often decision trees) to create a strong learner. It assigns weights to data points and adjusts them at each iteration based on whether the point was classified correctly.
 - **Gradient Boosting Machines (GBM):** Another popular boosting technique that builds trees sequentially and corrects errors of previous models by fitting new trees to the residual errors.
4. **Main Goal:**
- **Reduce Bias:** Boosting aims to reduce bias and improve accuracy by focusing on hard-to-classify examples, which leads to reduced error in predictions.

Key Differences:

- **Independence vs. Sequential Learning:** Bagging trains multiple models independently in parallel, while Boosting trains models sequentially, with each model learning from the mistakes of its predecessor.
- **Weighting:** Bagging gives equal weight to each model's prediction, while Boosting assigns higher weights to the instances that were misclassified by previous models.
- **Variance vs. Bias:** Bagging reduces variance and helps to prevent overfitting, while Boosting reduces bias and helps to improve accuracy.

Ans8. In Random Forests, the out-of-bag (OOB) error is a way to estimate the performance of the model without the need for a separate validation set. Here's an explanation of what out-of-bag error is and how it is calculated:

Out-of-Bag (OOB) Error in Random Forests:

1. **Bootstrap Sampling:**
 - Random Forests use an ensemble of decision trees, where each tree is trained on a bootstrap sample (random sample with replacement) of the training data.
 - Typically, about 2/3 of the data is used to train each tree, and the remaining 1/3 is not used (out-of-bag samples).
2. **Out-of-Bag (OOB) Samples:**
 - For each tree in the Random Forest, the samples that were not selected in the bootstrap sample serve as the out-of-bag samples for that tree.
 - These out-of-bag samples are not used in the training of that particular tree.
3. **Calculating OOB Error:**
 - After training the Random Forest, each tree can be tested using its corresponding out-of-bag samples.
 - The OOB error is then calculated as the prediction error of the model on the out-of-bag samples averaged over all trees in the forest.
4. **Advantages of OOB Error:**
 - **No Need for Cross-Validation:** OOB error provides an estimate of the model's performance without requiring a separate validation set or cross-validation.
 - **Efficient:** It leverages the data that was not used in the training of each individual tree, thus making efficient use of the available data.

- **Unbiased Estimate:** OOB error is an unbiased estimate of the true error (generalization error) of the Random Forest model.
5. **Usage:**
- The OOB error is useful for tuning hyperparameters of the Random Forest, such as the number of trees, maximum depth of trees, etc.
 - It can also be used to compare different versions of the model and to monitor the model's performance during training.

Ans9. K-fold cross-validation is a technique used to assess the performance of a machine learning model. It is a resampling procedure used to evaluate a model on a limited data sample, in situations where data is scarce.

Ans10. Hyperparameter tuning in machine learning refers to the process of choosing the optimal hyperparameters for a model. Hyperparameters are settings that are external to the model and whose values cannot be estimated from the data. They are typically set before the learning process begins. Examples of hyperparameters include the number of trees in a Random Forest, the learning rate of a Gradient Boosting Machine, the regularization parameter in a linear model, etc.

Importance of Hyperparameter Tuning:

1. **Model Performance:**
 - Hyperparameters significantly impact the performance of the model. Choosing the right hyperparameters can lead to a model that generalizes well to unseen data, while poor choices can result in overfitting or underfitting.
2. **Generalization:**
 - The goal of hyperparameter tuning is to improve the generalization ability of the model. A well-tuned model is expected to perform well on new, unseen data.
3. **Algorithm Behavior:**
 - Different hyperparameters can lead to significantly different models, and their optimal values can depend on the dataset and the specific problem being solved.

4. Process:

Ans11. If the learning rate in Gradient Descent is too large, it can lead to several issues:

1. **Divergence:**
 - The most critical issue with a large learning rate is that the gradient descent algorithm may diverge, rather than converge to the minimum of the loss function. Instead of minimizing the loss, the parameters may oscillate wildly or even grow without bound.
2. **Overshooting the Minimum:**

- With a large learning rate, each update to the parameters can be so large that the algorithm overshoots the minimum of the loss function. This can cause the parameters to bounce around the minimum, preventing convergence.
- 3. **Instability and Unpredictable Behavior:**
 - Large learning rates can lead to unstable behavior of the optimization process. The updates to the parameters may be so large and erratic that the optimization process becomes unpredictable.
- 4. **Poor Generalization:**
 - A large learning rate can cause the model to converge too quickly to a suboptimal solution that fits the training data well but does not generalize well to new, unseen data. This is because the model may overfit the training data.
- 5. **Convergence to Poor Local Minima:**
 - Gradient Descent with a large learning rate can cause the algorithm to converge to poor local minima or saddle points rather than the global minimum of the loss function. This is because the algorithm may not have the chance to settle into the basin of the global minimum due to the large steps.
- 6. **Inefficient Learning:**
 - Even if the model converges, a large learning rate may require more iterations to reach convergence. This is because the steps are so large that they miss the optimal solution and require more adjustments.

Mitigation Strategies:

To address issues with a large learning rate, several strategies can be employed:

- **Learning Rate Decay:** Reduce the learning rate over time during training to allow the model to settle into the minimum of the loss function more effectively.
- **Use of Adaptive Learning Rates:** Techniques like AdaGrad, RMSProp, and Adam adjust the learning rate for each parameter based on the historical gradients, which can help mitigate the issues caused by a large learning rate.
- **Validation Set:** Monitor the performance of the model on a separate validation set during training to detect issues with overfitting or poor generalization caused by a large learning rate.
- **Grid Search or Random Search:** Perform a search over a range of learning rates to find an optimal value that balances convergence speed and stability.

Ans12. Logistic Regression is a linear classifier and is not suitable for classifying non-linear data directly. Here's why:

Linearity of Logistic Regression:

1. **Linear Decision Boundary:**
 - Logistic Regression models the probability of the default class (typically class 1) via a linear combination of the predictor variables. Mathematically, the decision boundary for logistic regression is a linear function of the input features.

- For a binary classification problem with a single feature x , the logistic regression model can be represented as: $\text{logit}(p) = \beta_0 + \beta_1 x$ where p is the probability of the default class, and β_0 and β_1 are the coefficients.
2. **Limitation to Linear Separability:**
- Logistic Regression assumes that the decision boundary is linear. This means it can only learn linear decision boundaries between classes.
 - If the relationship between the features and the target variable is non-linear, logistic regression will not be able to model it effectively.

Handling Non-Linear Data:

1. **Transformations:**
 - One approach to using logistic regression for non-linear data is to transform the input features to make the relationship more linear. This can involve feature engineering or using basis functions.
 - For instance, adding polynomial terms or interaction terms to the model can help capture non-linear relationships to some extent.
2. **Feature Engineering:**
 - Create new features that are transformations of the original features. For example, if the data is non-linear in the original space, you could transform it to a higher-dimensional space where it becomes linear.
3. **Kernel Methods:**
 - Another approach is to use kernel methods, where you map the input features into a higher-dimensional space using a kernel function (e.g., polynomial kernel, Gaussian RBF kernel), where the problem may become linearly separable.

Alternative Models for Non-Linear Data:

1. **Decision Trees and Ensemble Methods:**
 - Decision tree-based models (e.g., Random Forests, Gradient Boosting Machines) are capable of capturing non-linear relationships and interactions between features.
 - These models can automatically model complex decision boundaries without the need for explicit feature transformations.
2. **Support Vector Machines (SVM):**
 - SVMs with non-linear kernels (e.g., polynomial kernel, Gaussian RBF kernel) can model non-linear decision boundaries effectively.
3. **Neural Networks:**
 - Neural networks, especially deep learning models, can learn complex non-linear relationships in the data. They are highly flexible and can model intricate decision boundaries.

Ans13. AdaBoost (Adaptive Boosting) and Gradient Boosting are both ensemble learning techniques that improve the predictive performance of machine learning models. However, they

have different approaches and are suited for different types of problems. Here's a comparison between AdaBoost and Gradient Boosting:

AdaBoost (Adaptive Boosting):

1. **Type of Learners:**
 - **Weak Learners:** AdaBoost combines multiple weak learners (e.g., decision stumps - shallow decision trees with one split) to create a strong learner.
 - **Sequential Training:** The weak learners are trained sequentially, and each subsequent learner focuses on the examples that were previously misclassified.
2. **Training Process:**
 - **Weight Adjustment:** AdaBoost assigns higher weights to incorrectly predicted instances so that subsequent weak learners pay more attention to them.
 - **Weighted Voting:** Each weak learner contributes to the final prediction with a weight that depends on its accuracy.
3. **Model Building:**
 - **Exponential Decrease:** AdaBoost uses an exponential loss function to update weights and focuses on hard examples.
 - **Classification:** AdaBoost is primarily used for classification problems.
4. **Advantages:**
 - **Simple to Implement:** AdaBoost is relatively simple and easy to implement.
 - **Less Prone to Overfitting:** It is less prone to overfitting compared to complex models.
5. **Disadvantages:**
 - **Sensitive to Noisy Data:** AdaBoost can be sensitive to noisy data and outliers.
 - **Bias-Variance Trade-off:** It can suffer from a high bias if the weak learner is too simple.
6. **Example:**
 - A decision stump is a common weak learner used in AdaBoost.

Gradient Boosting:

1. **Type of Learners:**
 - **Decision Trees:** Gradient Boosting builds an ensemble of decision trees, typically deep trees, to make predictions.
 - **Sequential Training:** Trees are trained sequentially, and each tree corrects errors made by the previous tree.
2. **Training Process:**
 - **Gradient Descent:** Gradient Boosting uses gradient descent algorithm to minimize the loss function.
 - **Residuals:** It fits new models to the residuals (errors) of prior models.
3. **Model Building:**
 - **Continuous Improvement:** Gradient Boosting improves the model by minimizing the loss function in the gradient direction.
 - **Flexibility:** It is a more flexible algorithm and can handle various types of data and models.

4. **Advantages:**

- **Highly Accurate:** Gradient Boosting typically provides better accuracy than AdaBoost.
- **Handles Complex Data:** It can model complex non-linear relationships.

5. **Disadvantages:**

- **Complexity:** Gradient Boosting is more complex and computationally intensive compared to AdaBoost.
- **Prone to Overfitting:** It can overfit if the number of trees is too large or the learning rate is too high.

6. **Example:**

- Gradient Boosting implementations include XGBoost, LightGBM, and sklearn's GradientBoostingClassifier/GradientBoostingRegressor.

Summary:

- **Approach:** AdaBoost combines multiple weak learners, whereas Gradient Boosting builds an ensemble of strong learners (typically decision trees).
- **Training:** AdaBoost uses weight adjustments and focuses on misclassified examples, while Gradient Boosting uses gradient descent and fits new models to residuals.
- **Model Type:** AdaBoost primarily uses decision stumps, while Gradient Boosting uses decision trees.

Ans14. The bias-variance trade-off is a fundamental concept in supervised machine learning that helps us understand the behavior of machine learning models, particularly in the context of their predictive performance. It refers to the problem of simultaneously minimizing two sources of error that prevent supervised learning algorithms from generalizing beyond the training data: bias and variance.

Bias:

- **Definition:** Bias refers to the error introduced by approximating a real-world problem, which is typically complex, by a simplified model.
- **Characteristics:** A high bias model is one that makes strong assumptions about the form of the target function (e.g., linear assumptions for a non-linear relationship).
- **Impact:** High bias can cause an algorithm to miss relevant relations between features and target outputs (underfitting).
- **Example:** A linear regression model applied to a dataset with a non-linear relationship between the input and output variables.

Variance:

- **Definition:** Variance refers to the error introduced by the algorithm's sensitivity to small fluctuations in the training dataset.
- **Characteristics:** A high variance model is one that is overly sensitive to noise in the training data.

- **Impact:** High variance can cause an algorithm to model the random noise in the training data rather than the intended outputs (overfitting).
- **Example:** A high-degree polynomial regression model applied to a dataset with limited observations.

Trade-off:

- **Bias-Variance Trade-off:** The goal in supervised learning is to find the right balance between bias and variance that minimizes the total error.
- **Underfitting:** High bias and low variance models often lead to underfitting, where the model is too simple to capture the underlying patterns in the data.
- **Overfitting:** Low bias and high variance models often lead to overfitting, where the model learns both the signal and the noise in the training data, which results in poor generalization to new data.

Practical Considerations:

- **Model Complexity:** Increasing the complexity of a model typically reduces bias but increases variance.
- **Regularization:** Regularization techniques can be used to reduce variance at the cost of introducing a controlled amount of bias.
- **Cross-Validation:** Techniques such as cross-validation can help evaluate the performance of a model and find the right balance between bias and variance.

Ans15. here's a short description of each type of kernel used in Support Vector Machines (SVM):

1. Linear Kernel:

- **Description:** The linear kernel is the simplest kernel function. It computes the dot product of the input features directly.
- **Formula:** $K(x, x') = x^T x'$
- **Usage:** Suitable for linearly separable data or when there are a large number of features.

2. RBF (Radial Basis Function) Kernel:

- **Description:** The RBF kernel is a popular choice because of its flexibility. It considers the similarity between points in a high-dimensional space.
- **Formula:** $K(x, x') = \exp(-\gamma \|x - x'\|^2)$
- **Usage:** Suitable for non-linear data. It maps the data into a higher-dimensional space where it becomes linearly separable.

3. Polynomial Kernel:

- **Description:** The polynomial kernel is a generalization of the linear kernel. It can capture non-linear relationships between the features.
- **Formula:** $K(x, x') = (\gamma x^T x' + r)^d$
 - γ is a scaling factor,
 - r is a coefficient, and
 - d is the degree of the polynomial.
- **Usage:** Suitable for non-linear problems where the data can be separated by polynomial boundaries.

Summary:

- **Linear Kernel:** Simple and efficient for linearly separable data.
- **RBF Kernel:** Flexible and effective for non-linear problems, widely used in practice.
- **Polynomial Kernel:** Captures complex relationships in the data, useful for non-linear problems with polynomial boundaries.